

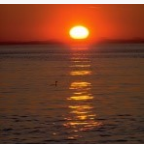


Chapter 4: Advanced SQL

Database System Concepts, 5th Ed.

©Silberschatz, Korth and Sudarshan

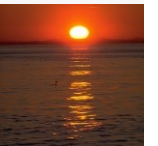
See www.db-book.com for conditions on re-use





Chapter 4: Advanced SQL

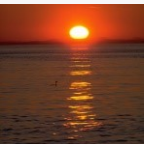
- SQL Data Types and Schemas
- Integrity Constraints
- Authorization
- Embedded SQL
- Dynamic SQL
- Functions and Procedural Constructs**
- Recursive Queries**
- Advanced SQL Features**





Built-in Data Types in SQL

- **date**: Dates, containing a (4 digit) year, month and date
 - Example: **date** '2005-7-27'
- **time**: Time of day, in hours, minutes and seconds.
 - Example: **time** '09:00:30' **time** '09:00:30.75'
- **timestamp**: date plus time of day
 - Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval**: period of time
 - Example: interval '1' day
 - Subtracting a date/time/timestamp value from another gives an interval value
 - Interval values can be added to date/time/timestamp values

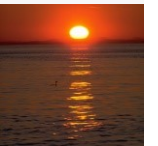




Build-in Data Types in SQL (Cont.)

- Can extract values of individual fields from date/time/timestamp
 - Example: **extract (year from r.starttime)**

- Can cast string types to date/time/timestamp
 - Example: **cast <string-valued-expression> as date**
 - Example: **cast <string-valued-expression> as time**





User-Defined Types

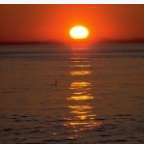
- **create type** construct in SQL creates user-defined type

create type *Dollars* **as numeric (12,2) final**

- **create domain** construct in SQL-92 creates user-defined domain types

create domain *person_name* **char(20) not null**

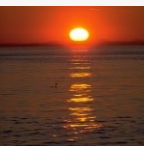
- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.





Domain Constraints

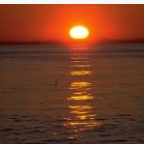
- **Domain constraints** are the most elementary form of integrity constraint. They test values inserted in the database, and test queries to ensure that the comparisons make sense.
- New domains can be created from existing data types
 - Example: **create domain Dollars numeric(12, 2)**
create domain Pounds numeric(12,2)
- We cannot assign or compare a value of type Dollars to a value of type Pounds.
 - However, we can convert type as below
(**cast r.A as Pounds**)
(Should also multiply by the dollar-to-pound conversion-rate)





Large-Object Types

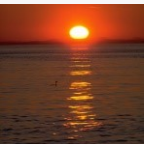
- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
 - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
 - **clob**: character large object -- object is a large collection of character data
 - When a query returns a large object, a pointer is returned rather than the large object itself.





Integrity Constraints

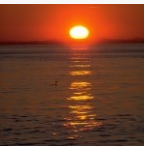
- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - A checking account must have a balance greater than \$10,000.00
 - A salary of a bank employee must be at least \$4.00 an hour
 - A customer must have a (non-null) phone number





Constraints on a Single Relation

- **not null**
- **primary key**
- **unique**
- **check (P)**, where P is a predicate

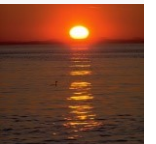




Not Null Constraint

- Declare *branch_name* for *branch* is **not null**
branch_name **char(15) not null**
- Declare the domain *Dollars* to be **not null**

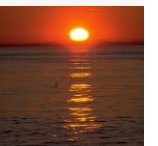
create domain *Dollars* numeric(12,2) not null





The Unique Constraint

- **unique** (A_1, A_2, \dots, A_m)
- The unique specification states that the attributes
 A_1, A_2, \dots, A_m
form a candidate key.
- Candidate keys are permitted to be null (in contrast to primary keys).



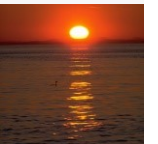


The check clause

- **check** (P), where P is a predicate

Example: Declare *branch_name* as the primary key for *branch* and ensure that the values of *assets* are non-negative.

```
create table branch
  (branch_name    char(15),
   branch_city   char(30),
   assets         integer,
   primary key (branch_name),
   check (assets >= 0))
```



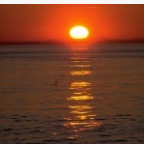


The check clause (Cont.)

- The **check** clause in SQL-92 permits domains to be restricted:
 - Use **check** clause to ensure that an `hourly_wage` domain allows only values greater than a specified value.

```
create domain hourly_wage numeric(5,2)  
           constraint value_test check(hourly_wage > = 4.00)
```

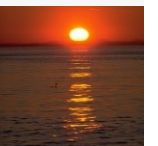
- The domain has a constraint that ensures that the `hourly_wage` is greater than 4.00
- The clause **constraint** *value_test* is optional; useful to indicate which constraint an update violated.





Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If “Perryridge” is a branch name appearing in one of the tuples in the *account* relation, then there exists a tuple in the *branch* relation for branch “Perryridge”.
- Primary and candidate keys and foreign keys can be specified as part of the SQL **create table** statement:
 - The **primary key** clause lists attributes that comprise the primary key.
 - The **unique key** clause lists attributes that comprise a candidate key.
 - The **foreign key** clause lists the attributes that comprise the foreign key and the name of the relation referenced by the foreign key. By default, a foreign key references the primary key attributes of the referenced table.





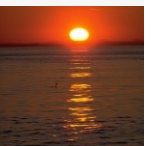
Referential Integrity in SQL – Example

create table *customer*

(*customer_name* **char**(20),
customer_street **char**(30),
customer_city **char**(30),
primary key (*customer_name*))

create table *branch*

(*branch_name* **char**(15),
branch_city **char**(30),
assets **numeric**(12,2),
primary key (*branch_name*))





Referential Integrity in SQL – Example (Cont.)

create table *account*

(account_number **char**(10),

branch_name **char**(15),

balance **integer**,

primary key (*account_number*),

foreign key (*branch_name*) **references** *branch*)

create table *depositor*

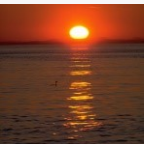
(customer_name **char**(20),

account_number **char**(10),

primary key (*customer_name*, *account_number*),

foreign key (*account_number*) **references** *account*,

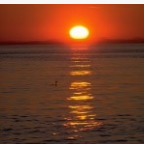
foreign key (*customer_name*) **references** *customer*)





Assertions

- An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
- An assertion in SQL takes the form
create assertion <assertion-name> **check** <predicate>
- When an assertion is made, the system tests it for validity, and tests it again on every update that may violate the assertion
 - This testing may introduce a significant amount of overhead; hence assertions should be used with great care.
- Asserting
for all X , $P(X)$
is achieved in a round-about fashion using
not exists X such that not $P(X)$

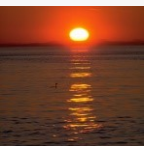




- **create assertion** <assertion-name> **check** <predicate>
- STUDENT(Roll no, name, Batch)

Create assertion Batch-size CHECK

```
(  
    ( SELECT BATCH , COUNT()  
      FROM STUDENT  
      GROUP BY BATCH HAVING COUNT <=20 )  
)
```

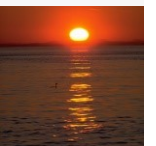




Assertion Example

- Every loan has at least one borrower who maintains an account with a minimum balance or \$1000.00

```
create assertion balance_constraint check  
  (not exists (  
    select *  
  from loan  
  where not exists (  
    select *  
    from borrower, depositor, account  
    where loan.loan_number = borrower.loan_number  
      and borrower.customer_name = depositor.customer_name  
      and depositor.account_number = account.account_number  
      and account.balance >= 1000)))
```

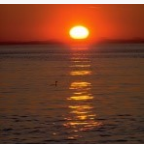




Assertion Example

- The sum of all loan amounts for each branch must be less than the sum of all account balances at the branch.

```
create assertion sum_constraint check  
  (not exists (select *  
    from branch  
    where (select sum(amount )  
      from loan  
        where loan.branch_name =  
          branch.branch_name )  
    >= (select sum (amount )  
      from account  
        where loan.branch_name =  
          branch.branch_name )))
```





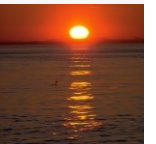
Authorization

Forms of authorization on parts of the database:

- **Read** - allows reading, but not modification of data.
- **Insert** - allows insertion of new data, but not modification of existing data.
- **Update** - allows modification, but not deletion of data.
- **Delete** - allows deletion of data.

Forms of authorization to modify the database schema (covered in Chapter 8):

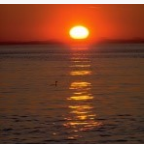
- **Index** - allows creation and deletion of indices.
- **Resources** - allows creation of new relations.
- **Alteration** - allows addition or deletion of attributes in a relation.
- **Drop** - allows deletion of relations.





Authorization Specification in SQL

- The **grant** statement is used to confer authorization
 grant <privilege list>
 on <relation name or view name> **to** <user list>
- <user list> is:
 - a user-id
 - **public**, which allows all valid users the privilege granted
 - A role (more on this in Chapter 8)
- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).



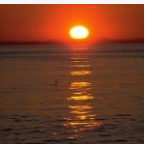


Privileges in SQL

- **select**: allows read access to relation, or the ability to query using the view
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *branch* relation:

grant select on *branch* to U_1 , U_2 , U_3

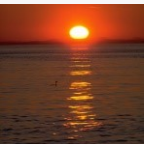
- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges
- more in Chapter 8





Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.
revoke <privilege list>
on <relation name or view name> **from** <user list>
- Example:
revoke select on branch from U_1, U_2, U_3
- All privileges that depend on the privilege being revoked are also revoked.
- <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.



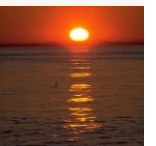


Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, Java, and Cobol.
- A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
- The basic form of these languages follows that of the System R embedding of SQL into PL/I.
- **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement > END_EXEC

Note: this varies by language (for example, the Java embedding uses
SQL { };)





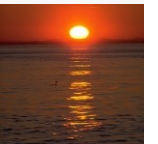
- Specify the query in SQL and declare a *cursor* for it

Eg: To find the customer names and their city

EXEC SQL

```
declare c cursor for  
select name, city  
from customer
```

END_EXEC

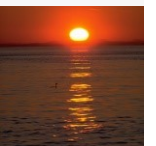




Example Query

- From within a host language, find the names and cities of customers with more than the variable **amount** dollars in some account.
- Specify the query in SQL and declare a *cursor* for it
EXEC SQL

```
declare c cursor for  
select depositor.customer_name, customer_city  
from depositor, customer, account  
where depositor.customer_name = customer.customer_name  
       and depositor.account_number = account.account_number  
       and account.balance > :amount  
  
END_EXEC
```





Embedded SQL (Cont.)

- The **open** statement causes the query to be evaluated

EXEC SQL **open** *c* END_EXEC

- The **fetch** statement causes the values of one tuple in the query result to be placed on host language variables.

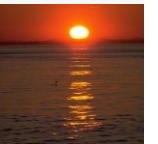
EXEC SQL **fetch** *c* **into** *:cn*, *:cc* END_EXEC

Repeated calls to **fetch** get successive tuples in the query result

- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available
- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

EXEC SQL **close** *c* END_EXEC

Note: above details vary with language. For example, the Java embedding defines Java iterators to step through result tuples.





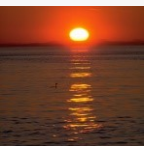
Updates Through Cursors

- Can update tuples fetched by cursor by declaring that the cursor is for update

```
declare c cursor for  
  select *  
  from account  
  where branch_name = 'Perryridge'  
for update
```

- To update tuple at the current location of cursor *c*

```
update account  
  set balance = balance + 100  
  where current of c
```





Dynamic SQL

- Allows programs to construct and submit SQL queries at run time.
- Example of the use of dynamic SQL from within a C program.

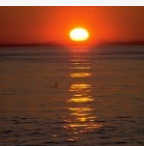
```
char * sqlprog = "update account  
                  set balance = balance * 1.05  
                  where account_number = ?"
```

```
EXEC SQL prepare dynprog from :sqlprog;
```

```
char account [10] = "A-101";
```

```
EXEC SQL execute dynprog using :account;
```

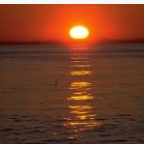
- The dynamic SQL program contains a ?, which is a place holder for a value that is provided when the SQL program is executed.





ODBC and JDBC

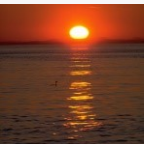
- API (application-program interface) for a program to interact with a database server
- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic
- JDBC (Java Database Connectivity) works with Java





ODBC

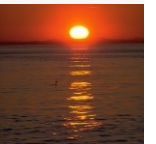
- Open DataBase Connectivity(ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - ▶ open a connection with a database,
 - ▶ send queries and updates,
 - ▶ get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC





ODBC (Cont.)

- Each database system supporting ODBC provides a "driver" library that must be linked with the client program.
- When client program makes an ODBC API call, the code in the library communicates with the server to carry out the requested action, and fetch results.
- ODBC program first allocates an SQL environment, then a database connection handle.
- Opens database connection using `SQLConnect()`. Parameters for `SQLConnect`:
 - connection handle,
 - the server to which to connect
 - the user identifier,
 - password
- Must also specify types of arguments:
 - `SQL_NTS` denotes previous argument is a null-terminated string.

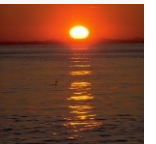




ODBC Code

```
■ int ODBCexample()
{
    RETCODE error;
    HENV  env;   /* environment */
    HDBC  conn; /* database connection */
    SQLAllocEnv(&env);
    SQLAllocConnect(env, &conn);
    SQLConnect(conn, "aura.bell-labs.com", SQL_NTS, "avi", SQL_NTS,
               "avipasswd", SQL_NTS);
    { .... Do actual work ... }

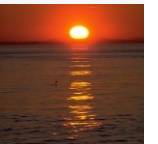
    SQLDisconnect(conn);
    SQLFreeConnect(conn);
    SQLFreeEnv(env);
}
```





ODBC Code (Cont.)

- Program sends SQL commands to the database by using `SQLExecDirect`
- Result tuples are fetched using `SQLFetch()`
- `SQLBindCol()` binds C language variables to attributes of the query result
 - When a tuple is fetched, its attribute values are automatically stored in corresponding C variables.
 - Arguments to `SQLBindCol()`
 - ▶ ODBC stmt variable, attribute position in query result
 - ▶ The type conversion from SQL to C.
 - ▶ The address of the variable.
 - ▶ For variable-length types like character arrays,
 - The maximum length of the variable
 - Location to store actual length when a tuple is fetched.
 - Note: A negative value returned for the length field indicates null value
- Good programming requires checking results of every function call for errors; we have omitted most checks for brevity.





ODBC Code (Cont.)

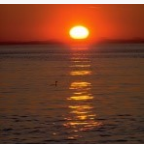
- Main body of program

```
char branchname[80];
float balance;
int lenOut1, lenOut2;
HSTMT stmt;

SQLAllocStmt(conn, &stmt);
char * sqlquery = "select branch_name, sum (balance)
                  from account
                  group by branch_name";

error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
if (error == SQL_SUCCESS) {
    SQLBindCol(stmt, 1, SQL_C_CHAR, branchname, 80,
&lenOut1);
    SQLBindCol(stmt, 2, SQL_C_FLOAT, &balance, 0,
&lenOut2);

    while (SQLFetch(stmt) >= SQL_SUCCESS) {
        printf (" %s %g\n", branchname, balance);
    }
}
SQLFreeStmt(stmt, SQL_DROP);
```





More ODBC Features

■ Prepared Statement

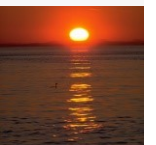
- SQL statement prepared: compiled at the database
- Can have placeholders: E.g. insert into account values(?,?,?)
- Repeatedly executed with actual values for the placeholders

■ Metadata features

- finding all the relations in the database and
- finding the names and types of columns of a query result or a relation in the database.

■ By default, each SQL statement is treated as a separate transaction that is committed automatically.

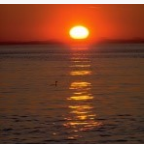
- Can turn off automatic commit on a connection
 - ▶ `SQLSetConnectOption(conn, SQL_AUTOCOMMIT, 0)`
- transactions must then be committed or rolled back explicitly by
 - ▶ `SQLTransact(conn, SQL_COMMIT)` or
 - ▶ `SQLTransact(conn, SQL_ROLLBACK)`





ODBC Conformance Levels

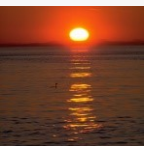
- Conformance levels specify subsets of the functionality defined by the standard.
 - Core
 - Level 1 requires support for metadata querying
 - Level 2 requires ability to send and retrieve arrays of parameter values and more detailed catalog information.
- SQL Call Level Interface (CLI) standard similar to ODBC interface, but with some minor differences.





JDBC

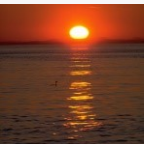
- **JDBC** is a Java API for communicating with database systems supporting SQL
- JDBC supports a variety of features for querying and updating data, and for retrieving query results
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes
- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors





JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
    try {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@aura.bell-labs.com:2000:bankdb", userid, passwd);
        Statement stmt = conn.createStatement();
        ... Do Actual Work ....
        stmt.close();
        conn.close();
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle);
    }
}
```





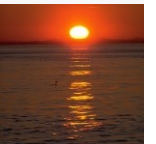
JDBC Code (Cont.)

- Update to database

```
try {  
    stmt.executeUpdate( "insert into account values  
                        ('A-9732', 'Perryridge', 1200)");  
} catch (SQLException sqle) {  
    System.out.println("Could not insert tuple. " + sqle);  
}
```

- Execute query and fetch and print results

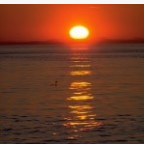
```
ResultSet rset = stmt.executeQuery( "select branch_name,  
                                     avg(balance)  
                                     from account  
                                     group by branch_name");  
  
while (rset.next()) {  
    System.out.println(  
        rset.getString("branch_name") + " " + rset.getFloat(2));  
}
```





JDBC Code Details

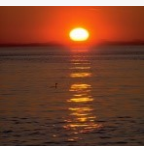
- Getting result fields:
 - **`rs.getString("branchname")` and `rs.getString(1)` equivalent if `branchname` is the first argument of select result.**
- Dealing with Null values
**`int a = rs.getInt("a");`
`if (rs.isNull()) Systems.out.println("Got null value");`**





Procedural Extensions and Stored Procedures

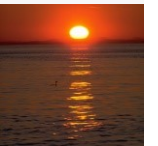
- SQL provides a **module** language
 - Permits definition of procedures in SQL, with if-then-else statements, for and while loops, etc.
 - more in Chapter 9
- Stored Procedures
 - Can store procedures in the database
 - then execute them using the **call** statement
 - permit external applications to operate on the database without knowing about internal details
- These features are covered in Chapter 9 (Object Relational Databases)





Functions and Procedures

- SQL:1999 supports functions and procedures
 - Functions/procedures can be written in SQL itself, or in an external programming language
 - Functions are particularly useful with specialized data types such as images and geometric objects
 - ▶ Example: functions to check if polygons overlap, or to compare images for similarity
 - Some database systems support **table-valued functions**, which can return a relation as a result
- SQL:1999 also supports a rich set of imperative constructs, including
 - Loops, if-then-else, assignment
- Many databases have proprietary procedural extensions to SQL that differ from SQL:1999





SQL Functions

- Define a function that, given the name of a customer, returns the count of the number of accounts owned by the customer.

```
create function account_count (customer_name varchar(20))  
returns integer  
begin  
    declare a_count integer;  
    select count ( * ) into a_count  
    from depositor  
    where depositor.customer_name = customer_name  
    return a_count;  
end
```

- Find the name and address of each customer that has more than one account.

```
select customer_name, customer_street, customer_city  
from customer  
where account_count (customer_name ) > 1
```

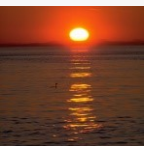




Table Functions

- SQL:2003 added functions that return a relation as a result
- Example: Return all accounts owned by a given customer

create function *accounts_of* (*customer_name* **char**(20)

returns table (*account_number* **char**(10),
branch_name **char**(15)
balance **numeric**(12,2))

return table

(**select** *account_number*, *branch_name*, *balance*
from *account A*
where exists (
 select *
 from *depositor D*
 where *D.customer_name* = *accounts_of.customer_name*
 and *D.account_number* = *A.account_number*))

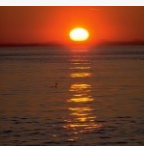
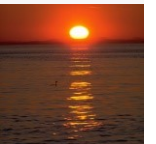




Table Functions (cont'd)

- Usage

```
select *  
from table (accounts_of ('Smith'))
```





SQL Procedures

- The *author_count* function could instead be written as procedure:

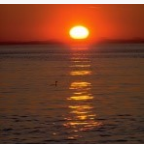
```
create procedure account_count_proc (in title varchar(20),  
                                     out a_count integer)  
  
begin  
    select count(author) into a_count  
    from depositor  
    where depositor.customer_name = account_count_proc.customer_name  
end
```

- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
declare a_count integer;  
call account_count_proc( 'Smith', a_count);
```

Procedures and functions can be invoked also from dynamic SQL

- SQL:1999 allows more than one function/procedure of the same name (called name **overloading**), as long as the number of arguments differ, or at least the types of the arguments differ





Procedural Constructs

- Compound statement: **begin ... end**,
 - May contain multiple SQL statements between **begin** and **end**.
 - Local variables can be declared within a compound statements
- **While** and **repeat** statements:

```
declare  $n$  integer default 0;
```

```
while  $n < 10$  do
```

```
    set  $n = n + 1$ 
```

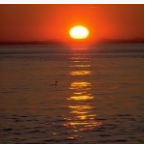
```
end while
```

```
repeat
```

```
    set  $n = n - 1$ 
```

```
until  $n = 0$ 
```

```
end repeat
```



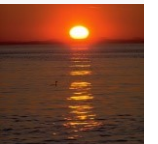


Procedural Constructs (Cont.)

■ For loop

- Permits iteration over all results of a query
- Example: find total of all balances at the Perryridge branch

```
declare n integer default 0;  
for r as  
    select balance from account  
    where branch_name = 'Perryridge'  
do  
    set n = n + r.balance  
end for
```





Procedural Constructs (cont.)

- Conditional statements (**if-then-else**)

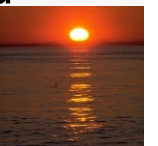
E.g. To find sum of balances for each of three categories of accounts (with balance <1000, >=1000 and <5000, >= 5000)

```
if r.balance < 1000
  then set l = l + r.balance
elseif r.balance < 5000
  then set m = m + r.balance
else set h = h + r.balance
end if
```

- SQL:1999 also supports a **case** statement similar to C case statement
- Signaling of exception conditions, and declaring handlers for exceptions

```
declare out_of_stock condition
declare exit handler for out_of_stock
begin
...
.. signal out-of-stock
end
```

- The handler here is **exit** -- causes enclosing **begin..end** to be exited
- Other actions possible on exception





External Language Functions/Procedures

- SQL:1999 permits the use of functions and procedures written in other languages such as C or C++
- Declaring external language procedures and functions

```
create procedure account_count_proc(in customer_name varchar(20),  
                                     out count integer)
```

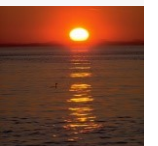
```
language C
```

```
external name ' /usr/avi/bin/account_count_proc'
```

```
create function account_count(customer_name varchar(20))  
returns integer
```

```
language C
```

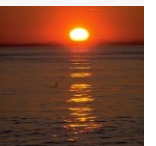
```
external name '/usr/avi/bin/author_count'
```





External Language Routines (Cont.)

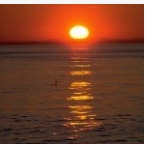
- Benefits of external language functions/procedures:
 - more efficient for many operations, and more expressive power
- Drawbacks
 - Code to implement function may need to be loaded into database system and executed in the database system's address space
 - ▶ risk of accidental corruption of database structures
 - ▶ security risk, allowing users access to unauthorized data
 - There are alternatives, which give good security at the cost of potentially worse performance
 - Direct execution in the database system's space is used when efficiency is more important than security





Security with External Language Routines

- To deal with security problems
 - Use **sandbox** techniques
 - ▶ that is use a safe language like Java, which cannot be used to access/damage other parts of the database code
 - Or, run external language functions/procedures in a separate process, with no access to the database process' memory
 - ▶ Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space



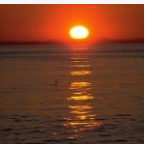


Recursion in SQL

- SQL:1999 permits recursive view definition
- Example: find all employee-manager pairs, where the employee reports to the manager directly or indirectly (that is manager's manager, manager's manager's manager, etc.)

```
with recursive empl (employee_name, manager_name ) as (  
    select employee_name, manager_name  
    from   manager  
    union  
    select manager.employee_name, empl.manager_name  
    from   manager, empl  
    where manager.manager_name = empl.employee_name)  
select *  
from   empl
```

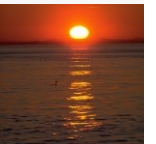
This example view, *empl*, is called the *transitive closure* of the *manager* relation





The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
 - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *manager* with itself
 - ▶ This can give only a fixed number of levels of managers
 - ▶ Given a program we can construct a database with a greater number of levels of managers on which the program will not work
- Computing transitive closure
 - The next slide shows a *manager* relation
 - Each step of the iterative process constructs an extended version of *empl* from its recursive definition.
 - The final result is called the *fixed point* of the recursive view definition.
- Recursive views are required to be *monotonic*. That is, if we add tuples to *manger* the view contains all of the tuples it contained before, plus possibly more

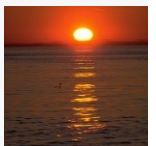




Example of Fixed-Point Computation

<i>employee_name</i>	<i>manager_name</i>
Alon	Barinsky
Barinsky	Estovar
Corbin	Duarte
Duarte	Jones
Estovar	Jones
Jones	Klinger
Rensal	Klinger

<i>Iteration number</i>	<i>Tuples in empl</i>
0	
1	(Duarte), (Estovar)
2	(Duarte), (Estovar), (Barinsky), (Corbin)
3	(Duarte), (Estovar), (Barinsky), (Corbin), (Alon)
4	(Duarte), (Estovar), (Barinsky), (Corbin), (Alon)

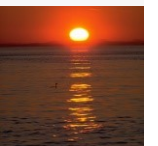




Advanced SQL Features**

- Create a table with the same schema as an existing table:
create table *temp_account* **like** *account*
- SQL:2003 allows subqueries to occur *anywhere* a value is required provided the subquery returns only one value. This applies to updates as well
- SQL:2003 allows subqueries in the **from** clause to access attributes of other relations in the **from** clause using the **lateral** construct:

```
select C.customer_name, num_accounts
from customer C,
      lateral (select count(*)
               from account A
               where A.customer_name = C.customer_name )
as this_customer (num_accounts )
```

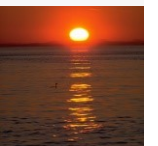




Advanced SQL Features (cont'd)

- Merge construct allows batch processing of updates.
- Example: relation *funds_received* (*account_number*, *amount*) has batch of deposits to be added to the proper account in the *account* relation

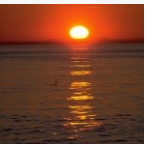
```
merge into account as A  
  using (select *  
          from funds_received as F )  
  on (A.account_number = F.account_number )  
  when matched then  
    update set balance = balance + F.amount
```

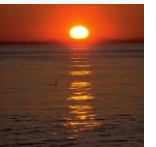




TRIGGERS

```
■ CREATE OR REPLACE TRIGGER t
BEFORE
INSERT OR UPDATE OF salary, department_id OR DELETE ON employees
BEGIN CASE
    WHEN INSERTING THEN DBMS_OUTPUT.PUT_LINE('Inserting');
    WHEN UPDATING('salary') THEN DBMS_OUTPUT.PUT_LINE('Updating salary');
    WHEN UPDATING('department_id') THEN DBMS_OUTPUT.PUT_LINE('Updating
        department ID');
    WHEN DELETING THEN DBMS_OUTPUT.PUT_LINE('Deleting'); END CASE;
END; /
```







creates a log table and a trigger that inserts a row in the log table after any **UPDATE** statement affects the **SALARY** column of the **EMPLOYEES** table, and then updates **EMPLOYEES.SALARY** and shows the log table.

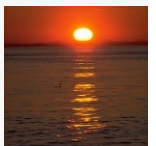
- CREATE TABLE Emp_log (Emp_id NUMBER, Log_date DATE, New_salary NUMBER, Action VARCHAR2(20));

- Create trigger that inserts row in log table after EMPLOYEES.SALARY is update

- CREATE OR REPLACE TRIGGER log_salary_increase
AFTER UPDATE OF salary ON employees

- FOR EACH ROW BEGIN INSERT INTO Emp_log (Emp_id, Log_date, New_salary, Action) VALUES (:NEW.employee_id, SYSDATE, :NEW.salary, 'New Salary');

- END;





- Update EMPLOYEES.SALARY:

UPDATE employees SET salary = salary + 1000.0 WHERE
Department_id = 20; Result:

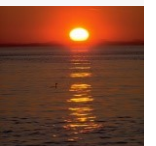
- 2 rows updated. Show log table:
- SELECT * FROM Emp_log;
- Result:

EMP_ID	LOG_DATE	NEW_SALARY	ACTION
--------	----------	------------	--------

201	28-APR-10	15049.13	New Salary
-----	-----------	----------	------------

202	28-APR-10	6945.75	New Salary
-----	-----------	---------	------------

2 rows selected.





End of Chapter

Database System Concepts, 5th Ed.

©Silberschatz, Korth and Sudarshan
See www.db-book.com for conditions on re-use

